AWESOME CYBER SECURITY HEAVEN — REALM DOSSIER · 2026 EDITION

# The AI Forge

## AI-Dependent Security Tools

| 35 | 608,574 | 82% | 27 |
|---|---|---|---|
| FREE TOOLS | COMMUNITY STARS | ON GITHUB ACTIVE | ACTIVE IN 2026 |

*35 tools · Part of Awesome Cyber Security Heaven — 1435 curated free tools across 14 realms · Data verified 2026-11-06 (re-verified monthly)*

## CURATOR'S NOTES

The AI-labour realm — security tools that put an LLM to work: red-team agents, pentest copilots, exercise generators. Every entry needs an LLM API key to function, which we mark as a feature, not a paywall — the tool is free, the brain is rented. We watch this realm closest of all: AI tooling churns faster than any other category on the list.

## MAINTENANCE SNAPSHOT (2026-11-06)

- **Active in the last 30 days:** 27 of 33 tracked GitHub repos (82%)
- **Quiet classics (no push in 12+ months, still listed):** none
- **Earliest entries:** PentestGPT (2023), promptfoo (2023), garak (2023)

## FLAGSHIP PICKS — THE HEAVYWEIGHTS

| TOOL | ★ STARS | CREATED | LAST PUSH | WHAT IT DOES |
|---|---|---|---|---|
| rea | 416.8k | 2026 | 2026-11-01 ✅ | Reverse engineer anything with agents, from app behavior down to native binaries. |
| strix | 73.2k | 2025 | 2026-10-30 ✅ | Open-source AI penetration testing tool to find and fix your app's vulnerabilities. |
| PentAGI | 32.9k | 2025 | 2026-10-22 ✅ | Fully autonomous AI Agents system capable of performing complex penetration testing tasks |
| promptfoo | 26.7k | 2023 | 2026-11-01 ✅ | Test your prompts, agents, and RAGs. Red teaming/pentesting/vulnerability scanning for AI. Compare performance... |
| PentestGPT | 16.2k | 2023 | 2026-07-14 | Automated Penetration Testing Agentic Framework Powered by Large Language Models |
| garak | 9.8k | 2023 | 2026-10-27 ✅ | the LLM vulnerability scanner |

*Full inventory — 35 curated tools · ranked by GitHub adoption · 2026-11-06 census · non-GitHub tools listed without stars*

| TOOL | ★ STARS | CREATED | LAST PUSH | LICENSE / TIER |
|---|---|---|---|---|
| rea | 416.8k | 2026 | 2026-11-01 ✅ | — open-source MCP coding agent that reproduces any web app's features from a live URL — the viral reference for agentic AI engineering, MCP-native and self-hosted. Used by researchers benchmarking reasoning-acting agents on security tasks; the open research agent for agentic testing. (MIT) |
| strix | 73.2k | 2025 | 2026-10-30 ✅ | — autonomous AI pentesting agent framework: sandboxed multi-agent teams, a 34-playbook vulnerability knowledge base, Caido MITM proxy integration, automated PoC validation and reporting. Used by security teams delegating scoped pentest workflows to supervised AI agents; the open-source agentic pentest framework with a large community. (Apache-2.0) |
| PentAGI | 32.9k | 2025 | 2026-10-22 ✅ | — fully autonomous AI-agent penetration testing: multi-agent workflows, browser control, reporting. Used by red teams experimenting with autonomous pentest agents; the open-source platform for agentic security testing. (MIT) |
| promptfoo | 26.7k | 2023 | 2026-11-01 ✅ | — LLM evaluation and red-teaming framework: prompt-injection and jailbreak test suites that run in CI. Used by AI engineers evaluating and red-teaming prompts in CI; the open-source testing framework for LLM quality and safety. (MIT) |

| | | | | |
|---|---|---|---|---|
| PentestGPT | 16.2k | 2023 | 2026-07-14 | — LLM-driven pentest copilot: parses recon output, plans next steps, keeps session context across test phases. Used by pentesters and learners getting guided next-step reasoning during engagements; the GPT-powered attack assistant. (MIT) |
| garak | 9.8k | 2023 | 2026-10-27 ✅ | — NVIDIA's LLM vulnerability scanner: hundreds of probes for injection, data leakage, hallucination and misuse. Used by AI security engineers probing LLMs for known vulnerability classes; NVIDIA's open-source LLM vulnerability scanner. (Apache-2.0) |
| T3MP3ST | 6.7k | 2026 | 2026-09-08 | — multi-agent offensive-security meta-framework that repurposes AI coding agents (Claude Code, Codex, local models via Ollama) as autonomous pentest brains, with a browser War Room UI and a verify-claims integrity layer; a reported 90% pass@1 score on the XBEN challenge suite. Used by red teams researching agentic kill chains reproducibly; the research harness for agent-driven offense. (AGPL-3.0, copyleft) |
| PyRIT | 4.7k | 2023 | 2026-11-01 ✅ | — Microsoft's AI red-teaming framework: automated adversarial probing of generative AI, with scoring and reporting. Used by AI red teams automating LLM risk probing; Microsoft's open-source toolkit for generative-AI security testing. (MIT) |
| RedAmon | 3.3k | 2025 | 2026-10-30 ✅ | — autonomous AI pentest framework chaining recon → exploit → triage → code-fix with human approval gates, a Neo4j attack-surface graph and bidirectional MCP support; its maintainers report a 97% black-box solve rate on the XBOW benchmark. Used by security teams researching supervised agentic testing; the approval-gated reference for AI-driven pentesting. (MIT) |
| deepteam | 3.1k | 2025 | 2026-10-23 ✅ | — red-teaming library for LLMs: simulated attacks with vulnerability scoring and reporting. Used by AI teams running red-teaming simulations on their LLMs; the open-source library for adversarial testing. (Apache-2.0) |
| Mantis | 2.7k | 2026 | 2026-10-30 ✅ | — Google's ADK-based toolkit that lets AI coding agents autonomously find, reproduce and patch vulnerabilities in source code — LLM-driven SAST with a severity-calibration rubric and anti-hallucination triage. Used by security engineers researching agentic vuln discovery; the flagship example of agent-driven security code review (needs paid LLM APIs). (Apache-2.0) |
| open-kritt | 2.4k | 2026 | 2026-10-29 ✅ | — Self-hosted AI vulnerability-research platform that orchestrates agents to find, validate, deduplicate and rank security issues in code. Used by security teams wanting AI-assisted code review without sending code to a vendor; the self-hosted alternative to commercial AI code-scanning SaaS.. (AGPL-3.0) |
| PentesterFlow Agent | 1.4k | 2026 | 2026-08-31 | — human-in-the-loop agentic pentest CLI: LLM-driven recon, enumeration, validation and reporting from the terminal, with evidence-gated anti-hallucination findings and full audit logging. Used by pentesters delegating scoped workflow steps to a supervised agent; the analyst-oversight-first agentic TUI. (Apache-2.0) |
| hackingBuddyGPT | 1.3k | 2023 | 2026-10-24 ✅ | — LLM pentesting research harness: autonomous Linux privilege-escalation experiments in ~50 lines of code. Used by researchers studying LLM agents in ethical hacking; the research framework for autonomous pentest agents. (MIT) |
| AttackGen | 1.3k | 2023 | 2026-10-24 ✅ | — LLM-powered, threat-informed exercise generation: realistic attack scenarios and mitigations for incident-response testing. Used by purple teams generating attack scenarios from threat intel; the open-source exercise-planning tool. (GPL-3.0) |
| ARTEX | 1.2k | 2026 | 2026-10-19 ✅ | — multi-agent LLM pentest orchestrator with human approval gates and asset graphing. Used by red teams running supervised multi-agent engagements; the open orchestrator for LLM-driven pentesting. (AGPL-3.0, copyleft) |
| DarkMoon | 1.2k | 2024 | 2026-10-27 ✅ | — autonomous AI pentest platform orchestrating 50 specialist agents across web, AD, cloud, Kubernetes and LLM endpoints (OWASP LLM Top 10), with a Privacy Gateway that tokenizes hosts and credentials locally before any LLM call. Used by teams researching full-scope agentic testing with evidence capture; dual-use — run only against authorized targets. (GPL-3.0) |
| Nebula | 1.1k | 2023 | 2026-10-31 ✅ | — cybersecurity harness that pairs an LLM with common pentest tools for guided exploitation workflows. Used by security researchers exploring LLM-assisted pentesting workflows; the open-source agent for offensive research. (BSD-2-Clause) |
| Cybermes | 987 | 2026 | 2026-10-31 ✅ | — autonomous offensive security, bug bounty and red-teaming agent framework with MCP integration and a built-in vulnerable lab. Used by offensive teams running autonomous agents; the Apache-licensed agentic framework with curated playbooks. (Apache-2.0) |

| | | | | |
|---|---|---|---|---|
| **Mitsu** | 537 | 2025 | 2026-10-24 ✅ | — chibi red-team operator: connect an AI provider, set the scope, and it drives recon through exploitation via 165+ MCP-connected tools. Used by red teamers running guided or autonomous authorized engagements; the MCP-native AI operator for red teams. (MIT) |
| **BugTraceAI** | 361 | 2026 | 2026-10-19 ✅ | — self-hosted agentic pentesting suite with evidence-first validation across an 8-repo ecosystem (CLI, WEB, API, Launcher). Used by security teams wanting self-hosted agentic scanning; the open alternative to SaaS AI pentest platforms. (Apache-2.0) |
| **Apex** | 322 | 2025 | 2026-10-30 ✅ | — autonomous offensive-security agents that run directly in the terminal, driving tools end-to-end on scoped targets. Used by pentesters delegating recon-to-exploit chains on authorized engagements; a terminal-native AI red-team operator. (Apache-2.0) |
| **Moxy** | 129 | 2026 | 2026-01-22 | — agentic MITM proxy with local-Ollama auto-testing built on mitmproxy. Used by web testers adding LLM-assisted traffic analysis; the open agentic layer for mitmproxy workflows. (Apache-2.0) |
| **PHANTOM** | 70 | 2026 | 2026-10-20 ✅ | — AI-powered pentesting command center with MCP and reverse-engineering integration. Used by pentesters orchestrating AI-assisted engagements; the open AI pentest cockpit. (MIT) |
| **WonderSuite Ai Bug Bounty** | 51 | 2026 | 2026-10-26 ✅ | — desktop-native AI offensive research engine with stealth browser and MCP agents. Used by bug bounty researchers automating recon-to-report workflows; the desktop AI engine for offensive research. (MIT) |
| **Xorcise** | 32 | 2026 | 2026-10-30 ✅ | — runs cyber-AI agents against prepared missions in isolated environments, with OpenTelemetry-based observability. Used by teams evaluating agent capability and safety before deployment; the open gym for offensive AI agents. (Apache-2.0) |
| **ReClass.NET-MCP** | 25 | 2025 | 2025-12-31 ⏳ | — MCP integration for ReClass.NET that exposes memory read/write, class/node management, and process inspection to AI coding assistants. Used by reverse engineers; an experimental bridge between LLM tooling and live memory analysis. |
| **VISTA** | 22 | 2025 | 2026-05-27 | — LLM-augmented Burp triage and workflow automation assistant. Used by Burp users triaging findings with LLM assistance; the open Burp copilot for triage. (MIT) |
| **Suijin** | 22 | 2026 | 2026-11-01 ✅ | — autonomous red-team framework on LangGraph with POC catalog and severity-honest reporting. (AGPL-3.0) |
| **bounty-operator** | 5 | 2026 | 2026-10-31 ✅ | — AI-powered bug bounty report triage that flags evidence gaps and duplicate risk before submission. Used by bug bounty hunters polishing reports before triage; the pre-submission report validator. (MIT) |
| **plonix** | 5 | 2026 | 2026-10-26 ✅ | — AI-native Rust MITM workbench with adaptive scope validation. Used by web security testers wanting an AI-native proxy; the Rust MITM workbench with LLM integration. (Apache-2.0) |
| **RedPilot** | 5 | 2026 | 2026-11-01 ✅ | — AI-orchestrated recon and scan pipeline chaining nuclei, nmap, ffuf, subfinder and httpx for authorized testing. Used by pentesters automating multi-tool recon under one AI orchestrator; the LLM-driven pipeline glue for standard recon stacks. (MIT) |
| **Hekate** | 1 | 2026 | 2026-10-20 ✅ | — Burp Suite passive extension with local-LLM triage and auto-validation to cut scanner false positives. Used by Burp users filtering scanner noise with local models; the open FP-reduction companion for Burp. (GPL-3.0) |
| **mobile-security-mcp** | — | — | — | — MCP server that gives AI agents automated access to Android reverse-engineering tools like apktool and jadx. Used by mobile security researchers; a novel AI-driven analysis bridge. |
| **Frida Game Hacking MCP** | — | — | — | — Model Context Protocol server that turns AI agents into game hacking assistants using Frida for memory operations, function hooking, and code patching. Used by AI-assisted reverse engineers; an experimental bridge between LLM agents and live process manipulation. |