The AI Forge

AI-Dependent Security Tools

20	433,698	85%	17
FREE TOOLS	COMBINED STARS	ACTIVE ON GITHUB	ACTIVE IN 2026

20 tools · ← Back to the main list · Data verified 2026-10-21 (re-verified monthly)

Curator's notes

The AI-labour realm — security tools that put an LLM to work: red-team agents, pentest copilots, exercise generators. Every entry needs an LLM API key to function, which we mark as a feature, not a paywall — the tool is free, the brain is rented. We watch this realm closest of all: AI tooling churns faster than any other category on the list.

Maintenance snapshot (2026-10-21)

- Active in the last 30 days: 9 of 10 tracked GitHub repos (90%)
- Youngest entries: PentestGPT (2023), promptfoo (2023), garak (2023)

Flagship picks

TOOL	★ STARS	CREATED	LAST PUSH	WHAT IT DOES
rea	337.3k	2026	2026-10-19 ✅	Reverse engineer anything with agents, from app behavior down to native binaries
PentAGI	30.3k	2025	2026-10-09 ✅	Fully autonomous AI Agents system capable of performing complex penetration testing tasks
promptfoo	26.0k	2023	2026-10-15 ✅	Test your prompts, agents, and RAGs. Red teaming/pentesting/vulnerability scanning for AI. Compare performance
PentestGPT	16.0k	2023	2026-07-14	Automated Penetration Testing Agentic Framework Powered by Large Language Models
garak	9.6k	2023	2026-10-15 ✅	the LLM vulnerability scanner
PyRIT	4.6k	2023	2026-10-15 ✅	The Python Risk Identification Tool for generative AI (PyRIT) is an open source framework built to empower security professionals and engineers to proactively i
deepteam	3.0k	2025	2026-10-01 ✅	DeepTeam is a framework to red team LLMs and AI agents.

Full inventory (20 tools)

TOOL	★ STARS	CREATED	LAST PUSH	LICENSE / TIER
PentAGI	30.3k	2025	2026-10-09 ✅	— fully autonomous AI-agent penetration testing: multi-agent workflows, browser control, reporting. (MIT)
promptfoo	26.0k	2023	2026-10-15 ✅	— LLM evaluation and red-teaming framework: prompt-injection and jailbreak test suites that run in CI. (MIT)
PentestGPT	16.0k	2023	2026-07-14	— LLM-driven pentest copilot: parses recon output, plans next steps, keeps session context across test phases. (MIT)
garak	9.6k	2023	2026-10-15 ✅	— NVIDIA's LLM vulnerability scanner: hundreds of probes for injection, data leakage, hallucination and misuse. (Apache-2.0)
PyRIT	4.6k	2023	2026-10-15 ✅	— Microsoft's AI red-teaming framework: automated adversarial probing of generative AI, with scoring and reporting. (MIT)
deepteam	3.0k	2025	2026-10-01 ✅	— red-teaming library for LLMs: simulated attacks with vulnerability scoring and reporting. (Apache-2.0)
ARTEX	1.6k	2026	2026-10-19 ✅	— multi-agent LLM pentest orchestrator with human approval gates and asset graphing — the open orchestrator for LLM-driven pentesting. (AGPL-3.0, copyleft)
hackingBuddyGPT	1.3k	2023	2026-10-01 ✅	— LLM pentesting research harness: autonomous Linux privilege-escalation experiments in ~50 lines of code. (MIT)
AttackGen	1.3k	2023	2026-10-13 ✅	— LLM-powered, threat-informed exercise generation: realistic attack scenarios and mitigations for incident-response testing. (GPL-3.0)
Nebula	1.1k	2023	2026-10-15 ✅	— cybersecurity harness that pairs an LLM with common pentest tools for guided exploitation workflows. (BSD-2-Clause)
Cybermes	953	2026	2026-10-21 ✅	— autonomous offensive security, bug bounty and red-teaming agent framework with MCP integration and a built-in vulnerable lab. (Apache-2.0)
BugTraceAI	365	2026	2026-10-19 ✅	— self-hosted agentic pentesting suite with evidence-first validation across an 8-repo ecosystem (CLI, WEB, API, Launcher). (Apache-2.0)
Moxy	128	2026	2026-01-22	— agentic MITM proxy with local-Ollama auto-testing built on mitmproxy — the open agentic layer for mitmproxy workflows. (Apache-2.0)
PHANTOM	70	2026	2026-10-20 ✅	— AI-powered pentesting command center with MCP and reverse-engineering integration — the open AI pentest cockpit. (MIT)
WonderSuite Ai Bug Bounty	51	2026	2026-10-19 ✅	— desktop-native AI offensive research engine with stealth browser and MCP agents — the desktop AI engine for offensive research. (MIT)
VISTA	21	2025	2026-05-27	— LLM-augmented Burp triage and workflow automation assistant — the open Burp copilot for triage. (MIT)
plonix	5	2026	2026-10-21 ✅	— AI-native Rust MITM workbench with adaptive scope validation — the Rust MITM workbench with LLM integration. (Apache-2.0)
bounty-operator	4	2026	2026-10-21 ✅	— AI-powered bug bounty report triage that flags evidence gaps and duplicate risk before submission. (MIT)

TOOL	★ STARS	CREATED	LAST PUSH	LICENSE / TIER
Hekate	1	2026	2026-10-20 ✅	— Burp Suite passive extension with local-LLM triage and auto-validation to cut scanner false positives. (GPL-3.0)
---	---	---	---	---

This page is auto-audited from the GitHub API on every monthly verification pass. Stars, ages and push dates are re-snapshotted; the main list only changes when a tool dies or a better one takes its place. See the editorial standards.