AWESOME CYBER SECURITY HEAVEN · REALM DOSSIER · 2026 EDITION

The AI Forge

AI-Dependent Security Tools

27	549,660	100%	27
FREE TOOLS	COMBINED STARS	ON GITHUB	ACTIVE IN 2026

27 tools · ← Back to the main list · Data verified **2026-10-22** (re-verified monthly)

CURATOR'S NOTES

The AI-labour realm — security tools that put an LLM to work: red-team agents, pentest copilots, exercise generators. Every entry needs an LLM API key to function, which we mark as a feature, not a paywall — the tool is free, the brain is rented. We watch this realm closest of all: AI tooling churns faster than any other category on the list.

MAINTENANCE SNAPSHOT (2026-10-22)

- Active in the last 30 days: 22 of 27 tracked GitHub repos (82%)
- Earliest entries: PentestGPT (2023), promptfoo (2023), garak (2023)

Part of Awesome Cyber Security Heaven — 496 curated free tools across 14 realms

FLAGSHIP PICKS — THE HEAVYWEIGHTS

TOOL	★ STARS	CREATED	LAST PUSH	WHAT IT DOES
rea	359.8k	2026	2026-10-22 ✅	Reverse engineer anything with agents, from app behavior down to native binaries.
strix	70.9k	2025	2026-10-22 ✅	Open-source AI penetration testing tool to find and fix your app's vulnerabilities.
PentAGI	31.8k	2025	2026-10-17 ✅	Fully autonomous AI Agents system capable of performing complex penetration testing tasks
promptfoo	26.3k	2023	2026-10-22 ✅	Test your prompts, agents, and RAGs. Red teaming/pentesting/vulnerability scanning for AI. Compare performance...
PentestGPT	16.1k	2023	2026-07-14	Automated Penetration Testing Agentic Framework Powered by Large Language Models
garak	9.6k	2023	2026-10-21 ✅	the LLM vulnerability scanner

FULL INVENTORY — 27 CURATED TOOLS

ranked by GitHub adoption · 2026-10-22 census · non-GitHub entries listed last

TOOL	★ STARS	CREATED	LAST PUSH	LICENSE / TIER
rea	359.8k	2026	2026-10-22 ✅	— open-source MCP coding agent that reproduces any web app's features from a live URL — the viral reference for agentic AI engineering, MCP-native and self-hosted. Used by researchers benchmarking reasoning-acting agents on security tasks; the open research agent for agentic testing. (MIT)

TOOL	★ STARS	CREATED	LAST PUSH	LICENSE / TIER
strix	70.9k	2025	2026-10-22 ✅	— autonomous AI pentesting agent framework: sandboxed multi-agent teams, a 34-playbook vulnerability knowledge base, Caido MITM proxy integration, automated PoC validation and reporting. Used by security teams delegating scoped pentest workflows to supervised AI agents; the open-source agentic pentest framework with a large community. (Apache-2.0)
PentAGI	31.8k	2025	2026-10-17 ✅	— fully autonomous AI-agent penetration testing: multi-agent workflows, browser control, reporting. Used by red teams experimenting with autonomous pentest agents; the open-source platform for agentic security testing. (MIT)
promptfoo	26.3k	2023	2026-10-22 ✅	— LLM evaluation and red-teaming framework: prompt-injection and jailbreak test suites that run in CI. Used by AI engineers evaluating and red-teaming prompts in CI; the open-source testing framework for LLM quality and safety. (MIT)
PentestGPT	16.1k	2023	2026-07-14	— LLM-driven pentest copilot: parses recon output, plans next steps, keeps session context across test phases. Used by pentesters and learners getting guided next-step reasoning during engagements; the GPT-powered attack assistant. (MIT)
garak	9.6k	2023	2026-10-21 ✅	— NVIDIA's LLM vulnerability scanner: hundreds of probes for injection, data leakage, hallucination and misuse. Used by AI security engineers probing LLMs for known vulnerability classes; NVIDIA's open-source LLM vulnerability scanner. (Apache-2.0)
T3MP3ST	6.6k	2026	2026-09-08	— multi-agent offensive-security meta-framework that repurposes AI coding agents (Claude Code, Codex, local models via Ollama) as autonomous pentest brains, with a browser War Room UI and a verify-claims integrity layer; a reported 90% pass@1 score on the XBEN challenge suite. Used by red teams researching agentic kill chains reproducibly; the research harness for agent-driven offense. (AGPL-3.0, copyleft)
PyRIT	4.7k	2023	2026-10-22 ✅	— Microsoft's AI red-teaming framework: automated adversarial probing of generative AI, with scoring and reporting. Used by AI red teams automating LLM risk probing; Microsoft's open-source toolkit for generative-AI security testing. (MIT)
RedAmon	3.2k	2025	2026-10-22 ✅	— autonomous AI pentest framework chaining recon → exploit → triage → code-fix with human approval gates, a Neo4j attack-surface graph and bidirectional MCP support; its maintainers report a 97% black-box solve rate on the XBOW benchmark. Used by security teams researching supervised agentic testing; the approval-gated reference for AI-driven pentesting. (MIT)
deepteam	3k	2025	2026-10-20 ✅	— red-teaming library for LLMs: simulated attacks with vulnerability scoring and reporting. Used by AI teams running red-teaming simulations on their LLMs; the open-source library for adversarial testing. (Apache-2.0)
Mantis	2.5k	2026	2026-10-20 ✅	— Google's ADK-based toolkit that lets AI coding agents autonomously find, reproduce and patch vulnerabilities in source code — LLM-driven SAST with a severity-calibration rubric and anti-hallucination triage. Used by security engineers researching agentic vuln discovery; the flagship example of agent-driven security code review (needs paid LLM APIs). (Apache-2.0)
ARTEX	1.6k	2026	2026-10-19 ✅	— multi-agent LLM pentest orchestrator with human approval gates and asset graphing. Used by red teams running supervised multi-agent engagements; the open orchestrator for LLM-driven pentesting. (AGPL-3.0, copyleft)
PentesterFlow Agent	1.4k	2026	2026-08-31	— human-in-the-loop agentic pentest CLI: LLM-driven recon, enumeration, validation and reporting from the terminal, with evidence-gated anti-hallucination findings and full audit logging. Used by pentesters delegating scoped workflow steps to a supervised agent; the analyst-oversight-first agentic TUI. (Apache-2.0)

TOOL	★ STARS	CREATED	LAST PUSH	LICENSE / TIER
hackingBuddyGPT	1.3k	2023	2026-10-01 ✅	— LLM pentesting research harness: autonomous Linux privilege-escalation experiments in ~50 lines of code. Used by researchers studying LLM agents in ethical hacking; the research framework for autonomous pentest agents. (MIT)
AttackGen	1.3k	2023	2026-10-19 ✅	— LLM-powered, threat-informed exercise generation: realistic attack scenarios and mitigations for incident-response testing. Used by purple teams generating attack scenarios from threat intel; the open-source exercise-planning tool. (GPL-3.0)
DarkMoon	1.1k	2024	2026-10-22 ✅	— autonomous AI pentest platform orchestrating 50 specialist agents across web, AD, cloud, Kubernetes and LLM endpoints (OWASP LLM Top 10), with a Privacy Gateway that tokenizes hosts and credentials locally before any LLM call. Used by teams researching full-scope agentic testing with evidence capture; dual-use — run only against authorized targets. (GPL-3.0)
Nebula	1.1k	2023	2026-10-22 ✅	— cybersecurity harness that pairs an LLM with common pentest tools for guided exploitation workflows. Used by security researchers exploring LLM-assisted pentesting workflows; the open-source agent for offensive research. (BSD-2-Clause)
Cybermes	953	2026	2026-10-21 ✅	— autonomous offensive security, bug bounty and red-teaming agent framework with MCP integration and a built-in vulnerable lab. Used by offensive teams running autonomous agents; the Apache-licensed agentic framework with curated playbooks. (Apache-2.0)
BugTraceAI	366	2026	2026-10-19 ✅	— self-hosted agentic pentesting suite with evidence-first validation across an 8-repo ecosystem (CLI, WEB, API, Launcher). Used by security teams wanting self-hosted agentic scanning; the open alternative to SaaS AI pentest platforms. (Apache-2.0)
Moxy	128	2026	2026-01-22	— agentic MITM proxy with local-Ollama auto-testing built on mitmproxy. Used by web testers adding LLM-assisted traffic analysis; the open agentic layer for mitmproxy workflows. (Apache-2.0)
PHANTOM	70	2026	2026-10-20 ✅	— AI-powered pentesting command center with MCP and reverse-engineering integration. Used by pentesters orchestrating AI-assisted engagements; the open AI pentest cockpit. (MIT)
WonderSuite Ai Bug Bounty	51	2026	2026-10-19 ✅	— desktop-native AI offensive research engine with stealth browser and MCP agents. Used by bug bounty researchers automating recon-to-report workflows; the desktop AI engine for offensive research. (MIT)
VISTA	21	2025	2026-05-27	— LLM-augmented Burp triage and workflow automation assistant. Used by Burp users triaging findings with LLM assistance; the open Burp copilot for triage. (MIT)
plonix	5	2026	2026-10-21 ✅	— AI-native Rust MITM workbench with adaptive scope validation. Used by web security testers wanting an AI-native proxy; the Rust MITM workbench with LLM integration. (Apache-2.0)
bounty-operator	4	2026	2026-10-21 ✅	— AI-powered bug bounty report triage that flags evidence gaps and duplicate risk before submission. Used by bug bounty hunters polishing reports before triage; the pre-submission report validator. (MIT)
RedPilot	4	2026	2026-10-22 ✅	— AI-orchestrated recon and scan pipeline chaining nuclei, nmap, ffuf, subfinder and httpx for authorized testing. Used by pentesters automating multi-tool recon under one AI orchestrator; the LLM-driven pipeline glue for standard recon stacks. (MIT)
Hekate	1	2026	2026-10-20 ✅	— Burp Suite passive extension with local-LLM triage and auto-validation to cut scanner false positives. Used by Burp users filtering scanner noise with local models; the open FP-reduction companion for Burp. (GPL-3.0)