AWESOME CYBER SECURITY HEAVEN — REALM DOSSIER · 2026 EDITION

The Oracle's Domain

AI Security & LLM Security Tools

24	88,585	62%	13
FREE TOOLS	COMMUNITY STARS	ON GITHUB ACTIVE	ACTIVE IN 2026

24 tools · Part of Awesome Cyber Security Heaven — 1481 curated free tools across 14 realms · Data verified 2026-11-14 (re-verified monthly)

CURATOR'S NOTES

The zero-AI AI-security realm — guardrails, model-file safety, signing and the shared risk taxonomies. Tools here defend AI surfaces without themselves needing an LLM key; the AI-dependent assistants and copilots live one realm over, in The AI Forge.

MAINTENANCE SNAPSHOT (2026-11-14)

- **Active in the last 30 days:** 12 of 20 tracked GitHub repos (60%)
- **Quiet classics (no push in 12+ months, still listed):** none
- **Earliest entries:** Awesome Threat Intelligence (2015), NeMo Guardrails (2023), modelscan (2023)

FLAGSHIP PICKS — THE HEAVYWEIGHTS

TOOL	★ STARS	CREATED	LAST PUSH	WHAT IT DOES
CL4R1T4S	51.7k	2025	2026-09-22	LEAKED SYSTEM PROMPTS FOR CHATGPT, CLAUDE, GEMINI, GROK, PERPLEXITY, CURSOR, LOVABLE, REPLIT, AND MORE! - AI S...
Awesome Threat Intelligence	10.8k	2015	2026-10-30 ✅	A curated list of Awesome Threat Intelligence resources
NeMo Guardrails	7.4k	2023	2026-10-30 ✅	NeMo Guardrails is an open-source toolkit for easily adding programmable guardrails to LLM-based conversationa...
modelscan	793	2023	2026-09-28	Protection against Model Serialization Attacks
HackAgent	532	2025	2026-10-31 ✅	HackAgent is an open-source security toolkit to detect vulnerabilities of your AI Agents
guardana	289	2026	2026-11-01 ✅	Open-source AI security verification for model artifacts, live endpoints, MCP servers, and recorded agent trac...

Full inventory — 24 curated tools · ranked by GitHub adoption · 2026-11-14 census · non-GitHub tools listed without stars

TOOL	★ STARS	CREATED	LAST PUSH	LICENSE / TIER
CL4R1T4S	51.7k	2025	2026-09-22	— the community archive of leaked and reverse-engineered system prompts from ChatGPT, Claude, Gemini, Grok, Cursor and other major models — provenance-tagged and continuously updated. Used by AI governance teams auditing prompt transparency and researchers studying instruction leakage; the reference corpus for prompt-disclosure audits. (AGPL-3.0, copyleft)
Awesome Threat Intelligence	10.8k	2015	2026-10-30 ✅	— the curated threat-intelligence tools and platform index. Used by CTI practitioners assembling sources and tools. (Apache-2.0, curated list)
Guardrails AI	7.6k	2023	2026-11-05 ✅	— Python framework adding input/output guards, validators, and structured-output controls to LLM applications. Used by AI engineers and security teams; a practical LLM safety and validation layer. (Apache-2.0)
NeMo Guardrails	7.4k	2023	2026-10-30 ✅	— programmable guardrail layer for LLM apps: input/output rails, dialog safety, topic control. Used by AI engineers adding safety rails to LLM apps; NVIDIA's open-source guardrail toolkit for production deployments.
PurpleLlama	4.5k	2023	2026-10-30 ✅	— Meta's open-source LLM security toolkit including Llama Guard classifiers, CodeShield, and CyberSecEval benchmarks. Used by AI developers and safety researchers; a comprehensive LLM safety evaluation suite. (NOASSERTION)
HarmBench	1.1k	2024	2024-08-16 ⏳	— ICML 2024 standardized evaluation framework for automated red-teaming and robust-refusal benchmarking of LLMs. Used by AI safety researchers; a rigorous methodology for measuring model harmfulness. (MIT)

LangKit	1.0k	2023	2024-11-22 ⏳	— LLM monitoring toolkit from WhyLabs extracting safety and security signals like jailbreak and prompt-injection similarity. Used by ML engineers and AI security teams; a practical runtime LLM guardrail library. (Apache-2.0)
EasyJailbreak	939	2024	2026-09-01 ⏳	— Framework for building and testing adversarial jailbreak prompts against LLMs with modular attack components. Used by AI security researchers; a flexible jailbreak experimentation toolkit. (GPL-3.0)
modelscan	793	2023	2026-09-28	— scan model files (pickle, safetensors) for malicious code before you load them — supply-chain defense for ML. Used by AI security teams scanning model files for pickle and format exploits; the open-source scanner for ML supply chains.
Fickling	676	2021	2026-11-05 ✅	— Pickle file analysis and decompilation tool detecting malicious pickle payloads through allowlist-based validation. Used by ML engineers and security researchers; a strong defense against pickle-based model attacks. (LGPL-3.0)
HackAgent	532	2025	2026-10-31 ✅	— red-team toolkit for AI models and agents: actively tests jailbreaks, prompt injection and tool misuse, then scores each reply with judge models for measured per-attack success rates. Used by AI security teams quantifying agent attack surface; the open-source AI red-teaming harness with built-in scoring. (Apache-2.0)
guardana	289	2026	2026-11-01 ✅	— AI-security verification platform: 58 deterministic rules across model artifacts, live LLM, MCP and A2A endpoints and OTel traces, mapped to OWASP LLM/Agentic/MCP, ATLAS and NIST. Used by AI security teams auditing model stacks deterministically; the pre-1.0 release-candidate platform worth tracking. (Apache-2.0)
model-transparency	249	2023	2026-10-26 ✅	— signing for ML models: verify exactly which model you are deploying. Used by AI governance teams documenting model provenance; the AI Alliance's open standard for model transparency labels. (Apache-2.0)
LLAMATOR	225	2024	2026-05-20 ⏳	— Red-teaming framework for chatbots and GenAI systems providing automated adversarial testing methodologies. Used by AI security testers; a structured conversational-AI safety evaluation tool. (NOASSERTION)
Argus	209	2026	2026-11-14 ✅	— AI-agent red-teaming framework: 274 black-box probes (OWASP LLM Top 10, MCP poisoning, browser-agent and x402 payment abuse) fired at a live agent through 15 target adapters, with rule-based verdicts and SARIF output. Used by AI security teams attacking their own agents first; the agent-focused red-team harness that also wraps 99 garak probes. (Apache-2.0)
Penthera	80	2026	2026-11-01 ✅	— offline rule-based security agent that audits app code, dependencies and AI coding tooling (Claude Code, Cursor, MCP servers) with deterministic fixes applied to a local branch — zero cloud calls. Used by teams hardening AI-assisted codebases with NIS2/CRA-ready evidence; the self-hosted AI-tooling auditor with SARIF output. (MIT)
weir	18	2026	2026-10-27 ✅	— leak tests for AI agents: deterministically finds every place an agent exposes credentials in telemetry, model calls or other services. Used by agent developers testing their own deployments; a focused open-source AI-agent leak auditor. (Apache-2.0)
FlareVM MCP Server	12	2025	2026-08-28 ⏳	— Model Context Protocol server enabling AI agents to drive an isolated, air-gapped FlareVM through 52 unified tools. Used by AI-assisted malware analysts; a bridge between LLM agents and analysis sandboxes. (MIT)
ZeroTrust	6	2025	2025-11-29 ⏳	— Privacy-first Chrome extension that analyzes website security locally with on-device WebGPU AI, scoring trust from HTTPS and phishing signals. Used by privacy-conscious users; a local AI security assessor. (MIT)
RedBlueAgent	1	2026	2026-11-14 ✅	— red-and-blue LLM security platform: PTES-driven pentest agent plus a standalone offline prompt-injection detector (AUROC 0.998) and a ContextGuard fusion layer that quarantines indirect prompt injection in tool output. Used by AI security teams hardening their own agents against LLM01 injection; the red-blue platform with a measured 100%-to-0% hijack reduction. (MIT)
OWASP Top 10 for LLM Applicatio…	—	—	—	— the standard risk taxonomy for LLM apps; the shared language for AI threat modeling. Used by AppSec teams auditing LLM systems; the industry-standard risk list everyone maps AI security to.
AgentHound	456	2026	2026-11-01 ✅	— offensive security framework for AI agent infrastructure (MCP/A2A/LiteLLM/Ollama/Qdrant/MLflow/Jupyter): differential-proof collector + queryable attack graph; DEF CON 34 featured. (Apache-2.0)
SmartScanner	—	—	—	— AI-driven web vulnerability assessment tool that automates security scanning for websites, APIs, and web applications. Used by web security testers; a commercial scanner with a free tier for automated checks. (free tier)

Meta Llama Prompt Guard 2	—	—	—	— Meta's prompt-injection classifier on HuggingFace, fine-tuned from mDeBERTa with adversarial-resistant tokenization. Used by LLM application developers and AI security teams; a lightweight prompt-injection detection model. (free, proprietary)